



In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

In [2]:

```
from sklearn.datasets import load_boston
```

In [3]:

```
boston = load_boston()
```

In [4]:

```
boston.keys()
```

Out[4]:

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

In [5]:

```
print(boston.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of black people by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> (<http://archive.ics.uci.edu/ml/machine-learning-databases/housing/>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics

...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

```
.. topic:: References
```

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [6]:

```
print(boston.data)
```

```
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

In [7]:

```
print(boston.target)
```

```
[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.  18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21.  12.7 14.5 13.2 13.1 13.5 18.9 20.  21.  24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20.  16.6 14.4 19.4 19.7 20.5 25.  23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16.  22.2 25.  33.  23.5 19.4 22.  17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20.  20.8 21.2 20.3 28.  23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22.  22.9 25.  20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22.  20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18.  14.3 19.2 19.6 23.  18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14.  14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17.  15.6 13.1 41.3 24.3 23.3 27.  50.  50.  50.  22.7 25.  50.  23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50.  32.  29.8 34.9 37.  30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.  22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25.  23.3 28.7 21.5 23.  26.7 21.7 27.5 30.1
 44.8 50.  37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.  24.  25.1 31.5
 23.7 23.3 22.  20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44.  50.  36.  30.1 33.8 43.1 48.8 31.  36.5 22.8
 30.7 50.  43.5 20.7 21.1 25.2 24.4 35.2 32.4 32.  33.2 33.1 29.1 35.1
 45.4 35.4 46.  50.  32.2 22.  20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29.  24.8 22.  26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21.  23.8 23.1
 20.4 18.5 25.  24.6 23.  22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19.  18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25.  19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50.  50.  50.  50.  50.  13.8 13.8 15.  13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3  8.8  7.2 10.5  7.4 10.2 11.5 15.1 23.2
  9.7 13.8 12.7 13.1 12.5  8.5  5.  6.3  5.6  7.2 12.1  8.3  8.5  5.
 11.9 27.9 17.2 27.5 15.  17.2 17.9 16.3  7.  7.2  7.5 10.4  8.8  8.4
 16.7 14.2 20.8 13.4 11.7  8.3 10.2 10.9 11.  9.5 14.5 14.1 16.1 14.3
 11.7 13.4  9.6  8.7  8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13.  13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20.  16.4 17.7
 19.5 20.2 21.4 19.9 19.  19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12.  14.6 21.4 23.  23.7 25.  21.8 20.6 21.2 19.1 20.6 15.2  7.
  8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22.  11.9]
```

In [8]:

```
print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

In [9]:

```
## Lets prepare the dataframe
```

In [10]:

```
dataset=pd.DataFrame(boston.data,columns=boston.feature_names)
dataset.head()
```

Out[10]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	1
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	1

In [11]:

```
dataset['Price']=boston.target
```

In [12]:

```
dataset.head()
```

Out[12]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	1
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	1

In [13]:

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   CRIM        506 non-null   float64
 1   ZN          506 non-null   float64
 2   INDUS       506 non-null   float64
 3   CHAS        506 non-null   float64
 4   NOX         506 non-null   float64
 5   RM          506 non-null   float64
 6   AGE         506 non-null   float64
 7   DIS         506 non-null   float64
 8   RAD         506 non-null   float64
 9   TAX         506 non-null   float64
10  PTRATIO     506 non-null   float64
11  B           506 non-null   float64
12  LSTAT       506 non-null   float64
13  Price       506 non-null   float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [14]:

```
dataset.describe()
```

Out[14]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	1000000.0



In [15]:

```
## Check the missing values
dataset.isnull().sum()
```

Out[15]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
Price     0
dtype: int64
```

In [16]:

```
## EDA
dataset.corr()
```

Out[16]:

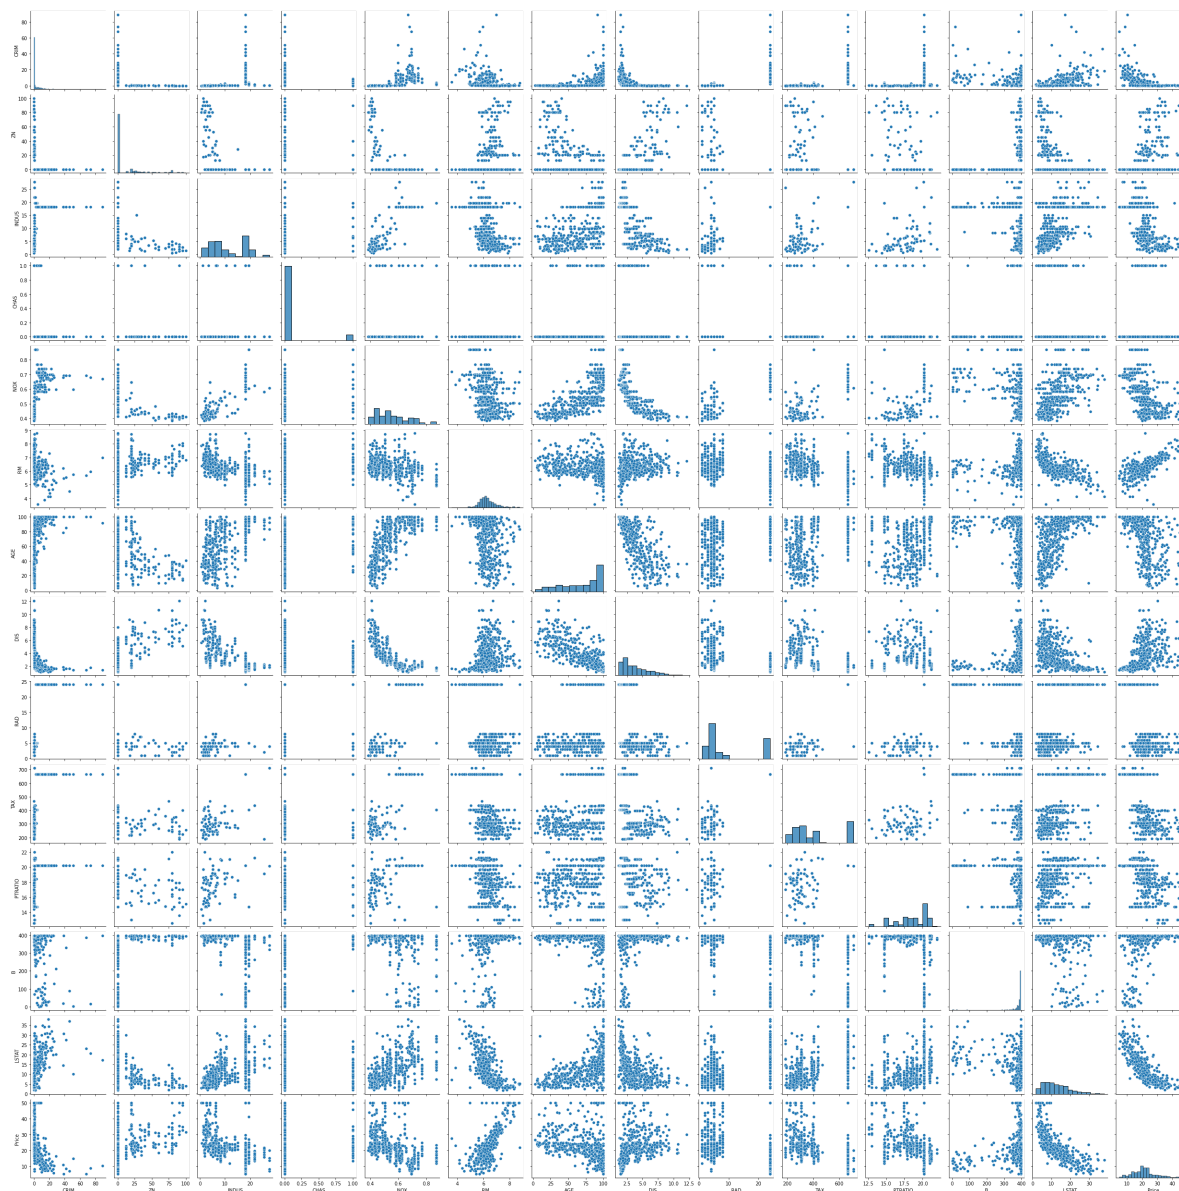
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929

In [17]:

```
import seaborn as sns
sns.pairplot(dataset)
```

Out[17]:

<seaborn.axisgrid.PairGrid at 0x180f666e5e0>




In [18]:

```
sns.set(rc={'figure.figsize':(10,8)})
sns.heatmap(dataset.corr(),annot=True)
```

Out[18]:

<AxesSubplot:>




In [19]:

```
plt.scatter(dataset['CRIM'],dataset['Price'])  
plt.xlabel("Crime Rate")  
plt.ylabel("Price")
```

Out[19]:

Text(0, 0.5, 'Price')




In [20]:

```
sns.set(rc={'figure.figsize':(8,6)})  
plt.scatter(dataset['RM'],dataset['Price'])
```

Out[20]:

<matplotlib.collections.PathCollection at 0x18081c8ce80>




In [21]:

```
sns.regplot(x='RM',y='Price',data=dataset)
```

Out[21]:

<AxesSubplot:xlabel='RM', ylabel='Price'>




In [22]:

```
sns.regplot(x="LSTAT",y="Price",data=dataset)
```

Out[22]:

<AxesSubplot:xlabel='LSTAT', ylabel='Price'>




In [23]:

```
sns.regplot(x='CRIM',y='Price',data=dataset)
```

Out[23]:

<AxesSubplot:xlabel='CRIM', ylabel='Price'>




In [24]:

```
sns.boxplot(dataset[ 'CRIM' ])
```

C:\Users\sekar\anaconda3\lib\site-packages\seaborn_decorators.py:36: Future Warning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[24]:

<AxesSubplot:xlabel='CRIM'>



In [25]:

dataset.head()

Out[25]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	5
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	5
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5

In [26]:

```
## Independent and Dependent Features
X = dataset.iloc[:, :-1]
y = dataset.iloc[:, -1]
```

In [27]:

X.head()

Out[27]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	5
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	5
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5

In [28]:

```
y
```

Out[28]:

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
```

Name: Price, Length: 506, dtype: float64

In [29]:

```
from sklearn.model_selection import train_test_split #test data validating the data
```

In [30]:

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.33,random_state=10)
```

In [31]:

```
X_train.shape
```

Out[31]:

```
(339, 13)
```

In [32]:

```
y_train.shape
```

Out[32]:

```
(339,)
```

In [33]:

```
X_test.shape
```

Out[33]:

```
(167, 13)
```

In [34]:

```
y_test.shape
```

Out[34]:

```
(167,)
```

In [35]:

```
y_test.shape
```

Out[35]:

```
(167,)
```

In [36]:

```
## Standardize or feature scaling the datasets
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

In [37]:

```
scaler
```

Out[37]:

```
StandardScaler()
```

In [38]:

```
X_train = scaler.fit_transform(X_train)
```

In [39]:

```
X_test=scaler.transform(X_test)
```

In [40]:

```
X_train
```

Out[40]:

```
array([[ -0.13641471, -0.47928013,  1.16787606, ..., -1.77731527,
         0.39261401,  2.36597873],
       [ -0.41777807, -0.47928013, -1.18043314, ..., -0.75987458,
         0.14721899, -0.54115799],
       [  1.31269177, -0.47928013,  0.95517731, ...,  0.76628645,
         0.19334986,  2.52100705],
       ...,
       [ -0.13520965, -0.47928013,  0.95517731, ...,  0.76628645,
         0.17012536,  0.06331026],
       [ -0.40281114, -0.47928013,  2.04022838, ...,  0.25756611,
         0.32166792,  0.27238516],
       [ -0.33104058,  0.34161649, -1.07552092, ..., -2.56351944,
         0.39993132, -0.34772815]])
```

In [41]:

X_test

Out[41]:

```
array([[ -0.41664568,  0.87519929, -1.33277144, ..., -0.06616502,
         0.41011193, -0.56391444],
       [ -0.42063267,  1.98340973, -1.22498491, ..., -1.36108953,
         0.41021798, -1.11860295],
       [ -0.41894074,  2.80430634, -1.16175014, ..., -1.12985301,
         0.44765291, -1.16980497],
       ...,
       [ -0.40804678,  1.36773726, -1.15169007, ..., -1.54607875,
         0.29854946, -1.18545003],
       [ -0.41098494, -0.47928013,  0.19779729, ...,  0.07257689,
         0.20119741, -0.13154186],
       [ -0.37856708, -0.47928013, -0.22328875, ..., -0.06616502,
         0.43482111, -0.5141347 ]])
```

Model Training

In [42]:

```
from sklearn.linear_model import LinearRegression
```

In [43]:

```
regression = LinearRegression()
```

In [44]:

regression

Out[44]:

LinearRegression()

In [45]:

```
regression.fit(X_train,y_train) #X train independet feature y train depfeaut
```

Out[45]:

LinearRegression()

In [46]:

```
##print the coefficents and the intercept
print(regression.coef_)
```

```
[ -1.29099218  1.60949999 -0.14031574  0.37201867 -1.76205329  2.22752218
   0.32268871 -3.31184248  2.70288107 -2.09005699 -1.7609799   1.25191514
  -3.83392028]
```

In [47]:

```
print(regression.intercept_)
```

22.077286135693214

In [48]:

```
## Prediction for the test data  
reg_pred=regression.predict(X_test)
```

In [49]:

```
reg_pred
```

Out[49]:

```
array([31.43849583, 31.98794389, 30.99895559, 22.31396689, 18.89492791,  
       16.21371128, 35.9881236 , 14.81264582, 25.04500847, 37.12806894,  
       21.49110158, 30.88757187, 28.05752881, 34.05600093, 33.75791114,  
       40.63880011, 24.24023412, 23.41351375, 25.54158122, 21.34135664,  
       32.71699711, 17.88341061, 25.49549436, 25.01006418, 32.54102925,  
       20.48979076, 19.48816948, 16.92733183, 38.38530857, 0.36265208,  
       32.42715816, 32.15306983, 26.10323665, 23.79611814, 20.67497128,  
       19.69393973, 3.50784614, 35.26259797, 27.04725425, 27.66164435,  
       34.35132103, 29.83057837, 18.40939436, 31.56953795, 17.91877807,  
       28.50042742, 19.49382421, 21.69553078, 38.0954563 , 16.44490081,  
       24.58507284, 19.67889486, 24.53954813, 34.30610423, 26.74699088,  
       34.87803562, 21.06219662, 19.87980936, 18.68725139, 24.71786624,  
       19.96344041, 23.56002479, 39.57630226, 42.81994338, 30.37060855,  
       17.03737245, 23.83719412, 3.2425022 , 31.5046382 , 28.63779884,  
       18.49288659, 27.14115768, 19.67125483, 25.34222917, 25.05430467,  
       10.29463949, 38.96369453, 8.26774249, 18.52214761, 30.34082002,  
       22.87681099, 20.96680268, 20.04604103, 28.73415756, 30.81726786,  
       28.23002473, 26.28588806, 31.59181918, 22.13093608, -6.48201197,  
       21.53000756, 19.90826887, 24.96686716, 23.44746617, 19.28521216,  
       18.75729874, 27.40013804, 22.17867402, 26.82972 , 23.39779064,  
       23.9260607 , 19.16632572, 21.09732823, 11.01452286, 13.7692535 ,  
       20.74596484, 23.54892211, 14.04445469, 28.88171403, 15.77611741,  
       15.25195598, 22.429474 , 26.60737213, 28.88742175, 24.29797261,  
       18.26839956, 16.26943281, 17.40100292, 15.53131616, 21.27868825,  
       33.78464602, 30.00899396, 21.16115702, 13.95560661, 16.18475215,  
       29.30998858, 13.1866784 , 22.08393725, 24.34499386, 31.86829501,  
       33.45923602, 5.90671516, 35.20153265, 24.17614831, 17.54200544,  
       24.25032915, 28.44671354, 34.50123773, 6.33164665, 1.93565618,  
       28.40727267, 12.56461105, 18.31045646, 19.71015745, 5.50105857,  
       14.51366874, 37.193992 , 25.81821367, 23.31632083, 26.43254504,  
       11.38255141, 20.46224115, 35.27645709, 20.57841598, 11.48799917,  
       16.23913171, 24.56511742, 10.53131603, 15.07115005, 25.98488217,  
       11.2136222 , 11.695686 , 19.40437966, 19.58768384, 32.43800883,  
       22.66170871, 25.68576052])
```

In [50]:


```
### Assumption of Linear Regression
```

In [51]:

```
plt.scatter(y_test,reg_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predicated Data')
```

Out[51]:

Text(0, 0.5, 'Test Predicated Data')



In [52]:

```
##residuals
residuals=y_test-reg_pred
```

In [53]:

```
residuals
```

Out[53]:


```
305    -3.038496
193    -0.887944
65     -7.498956
349     4.286033
151     0.705072
...
442    -1.004380
451    -4.387684
188    -2.638009
76     -2.661709
314    -1.885761
Name: Price, Length: 167, dtype: float64
```

In [54]:

```
sns.displot(residuals, kind= 'kde')
```

Out[54]:

<seaborn.axisgrid.FacetGrid at 0x18083078fd0>




In [55]:

```
## Scattte plot with preciction and residual  
## uniform distribution  
plt.scatter(reg_pred,residuals) #Doubt Model is good or not
```

Out[55]:

<matplotlib.collections.PathCollection at 0x18083179fa0>



In [56]:

```
## performance Metrics
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,reg_pred))
print(mean_absolute_error(y_test,reg_pred))
print(np.sqrt(mean_squared_error(y_test,reg_pred)))
```

```
27.100991709962482
3.520658529879791
5.205861284164464
```

R square and adjusted R square

In [57]:

```
from sklearn.metrics import r2_score
score = r2_score(y_test,reg_pred)
print(score)
```

```
0.7165219393967555
```

In [58]:

```
#Adjusted R2 square
#display adjusted R-squared
1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[58]:

```
0.6924355682343882
```

In [59]:

```
## Ridge
from sklearn.linear_model import Ridge
ridge=Ridge()
```

In [60]:

```
ridge.fit(X_train,y_train)
```

Out[60]:

```
Ridge()
```

In [61]:

```
ridge_pred = ridge.predict(X_test)
```


Assumptions of Ridge Regression

In [62]:

```
plt.scatter(y_test,ridge_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predicted Data')
```

Out[62]:

Text(0, 0.5, 'Test Predicted Data')



In [63]:

```
## Residual
ridge_residuals = y_test-ridge_pred
ridge_residuals
```

Out[63]:


```
305    -2.929516
193    -0.881807
65     -7.465240
349     4.148877
151     0.668281
...
442    -0.994515
451    -4.392072
188    -2.629490
76     -2.670984
314    -1.883764
Name: Price, Length: 167, dtype: float64
```

In [64]:

```
## Normal distributed expected  
sns.displot(ridge_residuals,kind='kde')
```

Out[64]:

<seaborn.axisgrid.FacetGrid at 0x1808316a280>




In [65]:

```
## Scatter plot with prediction and residual  
## If it is Uniform distribution, model is good  
plt.scatter(ridge_pred,ridge_residuals)
```

Out[65]:

<matplotlib.collections.PathCollection at 0x18085547be0>



In [67]:

```
## Performance Matrix  
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import mean_absolute_error  
print(mean_squared_error(y_test,ridge_pred))  
print(mean_absolute_error(y_test,ridge_pred))  
print(np.sqrt(mean_squared_error(y_test,ridge_pred)))
```

27.076490001440607

3.516104426348424

5.203507471066089

In [68]:

```
#R2 score  
from sklearn.metrics import r2_score  
r_score_ridge=r2_score(y_test,ridge_pred)  
print(r_score_ridge)
```

0.716778228793379

In [69]:

```
## Adjusted R square  
#display adjusted R-squared  
1-(1-r_score_ridge)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[69]:

0.6927136338542543

Lasso Regression

In [72]:

```
## Lasso  
from sklearn.linear_model import Lasso  
lasso=Lasso()
```

In [73]:

```
lasso.fit(X_train,y_train)
```

Out[73]:

Lasso()

In [74]:

```
#Lasso Prediction  
lasso_pred = lasso.predict(X_test)
```

In [75]:

```
lasso_pred
```

Out[75]:

```
array([25.64194382, 29.81425297, 27.94324255, 27.55256464, 20.99640298,
       18.74520609, 34.28217994, 15.93009427, 20.70883387, 34.07542731,
       19.90502439, 26.60490365, 24.07990755, 29.92866139, 29.22037693,
       36.40160499, 26.2514407 , 19.88117334, 23.967085 , 22.50869347,
       30.87428332, 18.78300957, 23.92041383, 25.68996484, 32.43275786,
       21.59346217, 20.77097939, 19.17145706, 34.09829244, 2.83421427,
       30.5699873 , 29.29565261, 26.85558827, 25.25346658, 19.26477827,
       19.73302762, 7.84289608, 29.77239449, 25.40207471, 25.60513357,
       32.3846261 , 26.89227407, 18.03007537, 29.36340326, 18.91119501,
       27.26813644, 20.46203931, 21.03622196, 34.39115891, 18.05973586,
       23.67935365, 18.6389767 , 22.52686697, 32.78082702, 26.03741902,
       30.39354515, 20.51475327, 20.94796259, 17.76992156, 24.71515119,
       21.39562999, 22.87363803, 36.66878913, 37.88636344, 28.19838095,
       17.67593653, 24.95639783, 5.16197744, 27.44103022, 33.73592095,
       19.49826488, 28.55803328, 19.79798303, 21.76346312, 24.39060267,
       10.72390653, 35.69760879, 8.87515321, 19.7342389 , 30.44923304,
       23.5731452 , 25.969039 , 21.53759864, 26.56759106, 29.08985079,
       27.11618087, 27.52241607, 31.13762774, 25.19009251, 1.294237 ,
       24.63082225, 20.35958514, 25.20780706, 25.3136586 , 19.78586427,
       20.81458041, 26.18578331, 20.65389528, 23.63766687, 22.07598695,
       24.48709112, 20.92346851, 21.99207783, 11.28215423, 15.85225507,
       22.37860011, 20.11691913, 17.8652717 , 27.92655648, 20.66334411,
       16.27505897, 21.25931319, 25.72062784, 25.92038934, 22.43890052,
       18.21105925, 13.73609426, 19.70671479, 19.16823222, 20.81634433,
       28.98351303, 31.19474714, 20.47647933, 14.48614909, 18.32476473,
       25.90552665, 15.20476004, 22.61771889, 24.11123534, 30.06699424,
       27.18076864, 10.97444888, 29.31613765, 29.11544566, 18.4851049 ,
       25.5516343 , 26.58527747, 29.79753531, 7.90378337, 13.9914479 ,
       27.53958439, 15.07581206, 19.1010104 , 20.2136539 , 9.56591918,
       12.35215313, 32.65105454, 27.04505972, 22.60389757, 23.38431025,
       18.14439345, 21.46916879, 33.35149732, 21.68370567, 11.686795 ,
       22.46129098, 24.462044 , 10.63445232, 17.62333713, 25.46239518,
       13.61074004, 15.25814185, 19.3752538 , 20.11939071, 29.36108994,
       22.66082095, 25.30493792])
```


Assumptions of Lasso Regression

In [76]:

```
plt.scatter(y_test,lasso_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predicted Data')
```

Out[76]:

Text(0, 0.5, 'Test Predicted Data')



In [77]:

```
## Residual
lasso_residuals = y_test-lasso_pred
lasso_residuals
```

Out[77]:


```
305    2.758056
193    1.285747
65     -4.443243
349    -0.952565
151    -1.396403
...
442    -0.975254
451    -4.919391
188     0.438910
76     -2.660821
314    -1.504938
Name: Price, Length: 167, dtype: float64
```

In [78]:

```
## Normal distribution expected
sns.displot(lasso_residuals,kind='kde')
```

Out[78]:

<seaborn.axisgrid.FacetGrid at 0x18087ac1640>




In [79]:

```
## Scatter plot with prediction and residual  
## If it is Uniform distribution , model is good  
plt.scatter(lasso_pred,lasso_residuals)
```

Out[79]:

<matplotlib.collections.PathCollection at 0x18088742b50>



In [80]:

```
## Performance Matrix  
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import mean_absolute_error  
print(mean_squared_error(y_test,lasso_pred))  
print(mean_absolute_error(y_test,lasso_pred))  
print(np.sqrt(mean_squared_error(y_test,lasso_pred)))
```

```
32.16822537607397  
3.9064325476573205  
5.671703921757021
```

In [81]:

```
#R2 score  
from sklearn.metrics import r2_score  
r_score_lasso=r2_score(y_test,lasso_pred)  
print(r_score_lasso)
```

0.6635183597615237

In [82]:

```
## Adjusted R square  
#display adjusted R-squared  
1-(1-r_score_lasso)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[82]:

0.6349284164732871

Elastic Net Regression

In [87]:

```
## Elastic Net  
from sklearn.linear_model import ElasticNet  
elasticnet = ElasticNet()
```

In [89]:

```
elasticnet.fit(X_train,y_train)
```

Out[89]:

ElasticNet()

In [91]:

```
## ElasticNet Prediction  
elasticnet_pred = elasticnet.predict(X_test)
```

In [92]:

```
elasticnet_pred
```

Out[92]:

```
array([26.0417533 , 29.72847396, 28.13249256, 27.33126697, 20.42880538,
       17.74088482, 31.34694254, 16.67485774, 22.66361605, 32.11606238,
       20.44062928, 27.05265082, 24.30388496, 29.10453835, 29.42032134,
       34.87404662, 25.31690008, 21.08018038, 24.04009667, 22.78241695,
       28.62957505, 18.35172223, 23.50225053, 24.94025282, 31.31440303,
       21.87551246, 22.30554751, 18.38033279, 33.5961939 ,  5.07350586,
       31.03524275, 28.19235387, 27.2862085 , 24.92462838, 19.28719449,
       20.2043877 ,  9.65913955, 29.64752478, 24.48773946, 25.34376165,
       30.68019641, 26.22751049, 18.01125345, 29.21052894, 20.61959202,
       27.27830384, 19.56149084, 19.72195809, 33.16071763, 19.16416141,
       23.05862027, 18.66118548, 22.77766754, 31.26962741, 25.0249516 ,
       29.94893114, 20.8407824 , 19.87498778, 18.27542547, 22.76517295,
       20.81723461, 22.76805785, 34.51940602, 36.11020157, 27.24493161,
       18.27047552, 24.17101249,  7.04406772, 26.85508847, 32.20329184,
       18.80351309, 27.46659498, 19.30788561, 20.8108208 , 23.90532445,
       14.22176442, 33.73545716, 10.78055539, 20.93445818, 29.78172162,
       23.99677889, 25.93581443, 22.61951728, 26.7428785 , 28.14408623,
       25.89892069, 26.67011775, 30.84900884, 24.58079972,  2.73998551,
       24.21010745, 20.83883219, 25.05619448, 24.60834531, 21.10174986,
       22.49049602, 26.06687733, 22.19194864, 23.68670917, 22.18041458,
       25.06905312, 20.26607354, 23.08760718, 10.49569995, 17.33370695,
       22.43253387, 21.15234101, 17.76115425, 26.58899485, 21.8834536 ,
       16.2603945 , 20.64765689, 25.09688902, 26.39218729, 22.53942481,
       18.04368235, 17.14638997, 20.67683019, 18.81134056, 19.82073711,
       26.78232308, 30.65570208, 20.62008364, 17.08768963, 19.18709774,
       25.65086934, 15.16746519, 21.83976175, 23.88029317, 29.16441724,
       27.25789963, 10.94687118, 29.43663882, 28.05096936, 18.27555631,
       24.83638128, 26.62936138, 30.15776119,  9.24328778, 10.68301626,
       27.04760052, 14.96584806, 18.73604079, 20.41851801,  9.92593546,
       14.91131492, 31.67918212, 27.61903571, 22.57544206, 24.72549427,
       15.35714158, 22.50035819, 31.86456612, 21.15334479, 12.67869784,
       22.65361671, 24.53650868, 11.9779066 , 18.36583206, 24.65573814,
       13.97917875, 14.8116782 , 18.92727223, 19.69541499, 28.7493839 ,
       22.61167089, 24.82104593])
```


Assumptions of ElasticNet Regression

In [93]:

```
plt.scatter(y_test,elasticnet_pred)  
plt.xlabel('Test Truth Data')  
plt.ylabel('Test Predicted Data')
```

Out[93]:

Text(0, 0.5, 'Test Predicted Data')



In [94]:

```
## Residual
elasticnet_residuals = y_test-elasticnet_pred
elasticnet_residuals
```

Out[94]:


```
305    2.358247
193    1.371526
65    -4.632493
349   -0.731267
151   -0.828805
...
442   -0.527272
451   -4.495415
188    1.050616
76    -2.611671
314   -1.021046
Name: Price, Length: 167, dtype: float64
```

In [95]:

```
## Normal distribution expected
sns.displot(elasticnet_residuals,kind='kde')
```

Out[95]:

<seaborn.axisgrid.FacetGrid at 0x18087ac74f0>




In [96]:

```
## Scatter plot with prediction and residual  
## If it is uniform distribution , model is good  
plt.scatter(elasticnet_pred,elasticnet_residuals)
```

Out[96]:

<matplotlib.collections.PathCollection at 0x180865df700>



In [98]:

```
## Performance Matrix  
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import mean_absolute_error  
print(mean_squared_error(y_test,elasticnet_pred))  
print(mean_absolute_error(y_test,elasticnet_pred))  
print(np.sqrt(mean_squared_error(y_test,elasticnet_pred)))
```

```
35.341543853934674  
4.035696708769101  
5.944875427957652
```


In [99]:

```
#R2 Score  
from sklearn.metrics import r2_score  
r_score_elasticnet = r2_score(y_test,elasticnet_pred)  
print(r_score_elasticnet)
```

0.6303252509112043

In [100]:

```
## Adjusted R sqaure  
# display adjusted R-squared  
1-(1-r_score_elasticnet)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
```

Out[100]:

0.5989149781128098

In [101]:

!pip install nbconvert[webpdf]

```

Requirement already satisfied: nbconvert[webpdf] in c:\users\sekar\anaconda3\lib\site-packages (6.1.0)
Requirement already satisfied: bleach in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (4.0.0)
Requirement already satisfied: traitlets>=5.0 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (5.1.0)
Requirement already satisfied: nbformat>=4.4 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (5.1.3)
Requirement already satisfied: nbclient<0.6.0,>=0.5.0 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.5.3)
Requirement already satisfied: jinja2>=2.4 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (2.11.3)
Requirement already satisfied: jupyterlab-pygments in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.1.2)
Requirement already satisfied: jupyter-core in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (4.8.1)
Requirement already satisfied: pygments>=2.4.1 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (2.10.0)
Requirement already satisfied: entrypoints>=0.2.2 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.3)
Requirement already satisfied: defusedxml in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.7.1)
Requirement already satisfied: mistune<2,>=0.8.1 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.8.4)
Requirement already satisfied: testpath in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (0.5.0)
Requirement already satisfied: pandocfilters>=1.4.1 in c:\users\sekar\anaconda3\lib\site-packages (from nbconvert[webpdf]) (1.4.3)
Collecting pypeteer==0.2.2
  Downloading pypeteer-0.2.2-py3-none-any.whl (145 kB)
Collecting websockets<9.0,>=8.1
  Downloading websockets-8.1.tar.gz (58 kB)
Requirement already satisfied: appdirs<2.0.0,>=1.4.3 in c:\users\sekar\anaconda3\lib\site-packages (from pypeteer==0.2.2->nbconvert[webpdf]) (1.4.4)
Requirement already satisfied: urllib3<2.0.0,>=1.25.8 in c:\users\sekar\anaconda3\lib\site-packages (from pypeteer==0.2.2->nbconvert[webpdf]) (1.26.7)
Requirement already satisfied: tqdm<5.0.0,>=4.42.1 in c:\users\sekar\anaconda3\lib\site-packages (from pypeteer==0.2.2->nbconvert[webpdf]) (4.62.3)
Collecting pyee<8.0.0,>=7.0.1
  Downloading pyee-7.0.4-py2.py3-none-any.whl (12 kB)
Requirement already satisfied: MarkupSafe>=0.23 in c:\users\sekar\anaconda3\lib\site-packages (from jinja2>=2.4->nbconvert[webpdf]) (2.0.1)
Requirement already satisfied: nest-asyncio in c:\users\sekar\anaconda3\lib\site-packages (from nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (1.5.1)
Requirement already satisfied: jupyter-client>=6.1.5 in c:\users\sekar\anaconda3\lib\site-packages (from nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (6.1.12)
Requirement already satisfied: async-generator in c:\users\sekar\anaconda3\lib\site-packages (from nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (1.10)
Requirement already satisfied: pyzmq>=13 in c:\users\sekar\anaconda3\lib\site-packages (from jupyter-client>=6.1.5->nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (22.2.1)

```

```
Requirement already satisfied: tornado>=4.1 in c:\users\sekar\anaconda3\lib\site-packages (from jupyter-client>=6.1.5->nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (6.1)
Requirement already satisfied: python-dateutil>=2.1 in c:\users\sekar\anaconda3\lib\site-packages (from jupyter-client>=6.1.5->nbclient<0.6.0,>=0.5.0->nbconvert[webpdf]) (2.8.2)
Requirement already satisfied: pywin32>=1.0 in c:\users\sekar\anaconda3\lib\site-packages (from jupyter-core->nbconvert[webpdf]) (228)
Requirement already satisfied: ipython-genutils in c:\users\sekar\anaconda3\lib\site-packages (from nbformat>=4.4->nbconvert[webpdf]) (0.2.0)
Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in c:\users\sekar\anaconda3\lib\site-packages (from nbformat>=4.4->nbconvert[webpdf]) (3.2.0)
Requirement already satisfied: pyrsistent>=0.14.0 in c:\users\sekar\anaconda3\lib\site-packages (from jsonschema!=2.5.0,>=2.4->nbformat>=4.4->nbconvert[webpdf]) (0.18.0)
Requirement already satisfied: setuptools in c:\users\sekar\anaconda3\lib\site-packages (from jsonschema!=2.5.0,>=2.4->nbformat>=4.4->nbconvert[webpdf]) (58.0.4)
Requirement already satisfied: six>=1.11.0 in c:\users\sekar\anaconda3\lib\site-packages (from jsonschema!=2.5.0,>=2.4->nbformat>=4.4->nbconvert[webpdf]) (1.16.0)
Requirement already satisfied: attrs>=17.4.0 in c:\users\sekar\anaconda3\lib\site-packages (from jsonschema!=2.5.0,>=2.4->nbformat>=4.4->nbconvert[webpdf]) (21.2.0)
Requirement already satisfied: colorama in c:\users\sekar\anaconda3\lib\site-packages (from tqdm<5.0.0,>=4.42.1->pyppeteer==0.2.2->nbconvert[webpdf]) (0.4.4)
Requirement already satisfied: webencodings in c:\users\sekar\anaconda3\lib\site-packages (from bleach->nbconvert[webpdf]) (0.5.1)
Requirement already satisfied: packaging in c:\users\sekar\anaconda3\lib\site-packages (from bleach->nbconvert[webpdf]) (21.0)
Requirement already satisfied: pyparsing>=2.0.2 in c:\users\sekar\anaconda3\lib\site-packages (from packaging->bleach->nbconvert[webpdf]) (3.0.4)
Building wheels for collected packages: websockets
  Building wheel for websockets (setup.py): started
  Building wheel for websockets (setup.py): finished with status 'done'
  Created wheel for websockets: filename=websockets-8.1-cp39-cp39-win_amd64.whl size=62758 sha256=816b4578c8e95e0d6f2b3a6b9e7080b64752f253a001ab39f959f01974efac18
  Stored in directory: c:\users\sekar\appdata\local\pip\cache\wheels\d8\b9\xa0\b97b211aeda2ebd6ac2e43fc300d308dbf1f9df520ed390cae
Successfully built websockets
Installing collected packages: websockets, pyee, pyppeteer
Successfully installed pyee-7.0.4 pyppeteer-0.2.2 websockets-8.1
```

In []: